

# mllabpractical

April 23, 2024

```
[7]: # Import necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer
from sklearn.ensemble import RandomForestClassifier, \
    GradientBoostingClassifier, VotingClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score, \
    f1_score, roc_auc_score, confusion_matrix

# Step 1: Data Preparation
# Load the dataset
df = pd.read_csv("hr_train.csv")

# Explore the dataset
print(df.head())
print(df.info())
print(df.describe())

# Perform data cleaning and preprocessing
# Handle missing values
# Let's fill missing values with the median for numerical columns
# For categorical features, let's fill missing values with the most frequent
    value
imputer = SimpleImputer(strategy='most_frequent')
df[categorical_features] = imputer.fit_transform(df[categorical_features])
numerical_features = df.select_dtypes(include=['int64', 'float64']).columns
categorical_features = df.select_dtypes(include=['object']).columns

imputer = SimpleImputer(strategy='median')
df[numerical_features] = imputer.fit_transform(df[numerical_features])
```

```

# For categorical features, let's fill missing values with the most frequent
↳ value
imputer = SimpleImputer(strategy='most_frequent')
df[categorical_features] = imputer.fit_transform(df[categorical_features])

# Handle categorical variables
# One-hot encode categorical features
encoder = OneHotEncoder(handle_unknown='ignore')
encoded_features = pd.DataFrame(encoder.fit_transform(df[categorical_features]).
↳ toarray(),
                                columns=encoder.
↳ get_feature_names_out(categorical_features))

# Concatenate encoded features with numerical features
df_encoded = pd.concat([df[numerical_features], encoded_features], axis=1)

# Step 2: Handling Imbalanced Data
# Investigate class imbalance
print(df['left'].value_counts())

# Implement resampling techniques
# Let's use SMOTE for oversampling
from imblearn.over_sampling import SMOTE

X = df_encoded.drop('left', axis=1)
y = df_encoded['left']

smote = SMOTE(random_state=42)
X_resampled, y_resampled = smote.fit_resample(X, y)

# Step 3: Feature Engineering
# No additional feature engineering needed for now

# Step 4: Feature Importance Analysis
# Train a random forest classifier to get feature importance
rf_classifier = RandomForestClassifier()
rf_classifier.fit(X_resampled, y_resampled)

# Get feature importances
feature_importances = rf_classifier.feature_importances_
sorted_idx = np.argsort(feature_importances)[::-1]

# Visualize feature importance
plt.figure(figsize=(10, 6))
sns.barplot(x=feature_importances[sorted_idx], y=X.columns[sorted_idx])
plt.title("Feature Importance")

```

```

plt.xlabel("Importance")
plt.ylabel("Features")
plt.show()

# Step 5: Model Building
# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_resampled, y_resampled,
    ↪test_size=0.2, random_state=42)

# Choose appropriate machine learning algorithms
# Let's use Logistic Regression, Random Forest, and Gradient Boosting
models = {
    "Logistic Regression": LogisticRegression(),
    "Random Forest": RandomForestClassifier(),
    "Gradient Boosting": GradientBoostingClassifier()
}

# Train multiple models
for name, model in models.items():
    model.fit(X_train, y_train)

# Step 6: Model Evaluation
# Evaluate the performance of each model
for name, model in models.items():
    y_pred = model.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    precision = precision_score(y_test, y_pred)
    recall = recall_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)
    roc_auc = roc_auc_score(y_test, y_pred)
    print(f"{name}:")
    print(f"Accuracy: {accuracy}")
    print(f"Precision: {precision}")
    print(f"Recall: {recall}")
    print(f"F1 Score: {f1}")
    print(f"ROC AUC Score: {roc_auc}")
    print()

# Step 7: Hyperparameter Tuning
# Let's perform hyperparameter tuning for Random Forest using GridSearchCV
param_grid = {
    "n_estimators": [100, 200, 300],
    "max_depth": [None, 10, 20],
    "min_samples_split": [2, 5, 10],
    "min_samples_leaf": [1, 2, 4],
    "bootstrap": [True, False]
}

```

```

rf = RandomForestClassifier()
grid_search = GridSearchCV(estimator=rf, param_grid=param_grid, cv=3,
    ↪scoring='accuracy')
grid_search.fit(X_train, y_train)

# Print best parameters
print("Best Parameters:")
print(grid_search.best_params_)

# Evaluate the tuned model
best_rf = grid_search.best_estimator_
y_pred = best_rf.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
roc_auc = roc_auc_score(y_test, y_pred)
print("Tuned Random Forest:")
print(f"Accuracy: {accuracy}")
print(f"Precision: {precision}")
print(f"Recall: {recall}")
print(f"F1 Score: {f1}")
print(f"ROC AUC Score: {roc_auc}")

# Step 8: Ensemble Methods
# Let's use Voting Classifier with Logistic Regression, Random Forest, and
    ↪Gradient Boosting
voting_classifier = VotingClassifier(
    estimators=[('lr', models['Logistic Regression']),
                ('rf', models['Random Forest']),
                ('gb', models['Gradient Boosting'])],
    voting='soft'
)

# Train the Voting Classifier
voting_classifier.fit(X_train, y_train)

# Evaluate the Voting Classifier
y_pred = voting_classifier.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
roc_auc = roc_auc_score(y_test, y_pred)
print("Voting Classifier:")
print(f"Accuracy: {accuracy}")

```

```

print(f"Precision: {precision}")
print(f"Recall: {recall}")
print(f"F1 Score: {f1}")
print(f"ROC AUC Score: {roc_auc}")

# Step 9: Model Interpretation
# Feature importance is already discussed above

# Plot misclassification error vs number of trees for Bagging and Boosting
↳ algorithms

# Let's use Random Forest and Gradient Boosting for this
train_errors = []
test_errors = []

# For Random Forest
for n_estimators in range(1, 101):
    rf_model = RandomForestClassifier(n_estimators=n_estimators)
    rf_model.fit(X_train, y_train)
    train_errors.append(1 - rf_model.score(X_train, y_train))
    test_errors.append(1 - rf_model.score(X_test, y_test))

plt.plot(range(1, 101), train_errors, label="Training Error")
plt.plot(range(1, 101), test_errors, label="Test Error")
plt.xlabel("Number of Trees")
plt.ylabel("Misclassification Error")
plt.title("Misclassification Error vs Number of Trees (Random Forest)")
plt.legend()
plt.show()

# For Gradient Boosting
train_errors = []
test_errors = []

for n_estimators in range(1, 101):
    gb_model = GradientBoostingClassifier(n_estimators=n_estimators)
    gb_model.fit(X_train, y_train)
    train_errors.append(1 - gb_model.score(X_train, y_train))
    test_errors.append(1 - gb_model.score(X_test, y_test))

plt.plot(range(1, 101), train_errors, label="Training Error")
plt.plot(range(1, 101), test_errors, label="Test Error")
plt.xlabel("Number of Trees")
plt.ylabel("Misclassification Error")
plt.title("Misclassification Error vs Number of Trees (Gradient Boosting)")
plt.legend()
plt.show()

```

	satisfaction_level	last_evaluation	number_project	average_monthly_hours	\
0	0.42	0.46	2	150	
1	0.66	0.77	2	171	
2	0.55	0.49	5	240	
3	0.22	0.88	4	213	
4	0.20	0.72	6	224	

	time_spend_company	Work_accident	left	promotion_last_5years	sales	\
0	3	0	1	0	sales	
1	2	0	0	0	technical	
2	3	0	0	0	technical	
3	3	1	0	0	technical	
4	4	0	1	0	technical	

salary

0	medium
1	medium
2	high
3	medium
4	medium

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 10499 entries, 0 to 10498

Data columns (total 10 columns):

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	satisfaction_level	10499 non-null	float64
1	last_evaluation	10499 non-null	float64
2	number_project	10499 non-null	int64
3	average_monthly_hours	10499 non-null	int64
4	time_spend_company	10499 non-null	int64
5	Work_accident	10499 non-null	int64
6	left	10499 non-null	int64
7	promotion_last_5years	10499 non-null	int64
8	sales	10499 non-null	object
9	salary	10499 non-null	object

dtypes: float64(2), int64(6), object(2)

memory usage: 820.4+ KB

None

	satisfaction_level	last_evaluation	number_project	\
count	10499.000000	10499.000000	10499.000000	
mean	0.612683	0.717131	3.808553	
std	0.248578	0.171483	1.230572	
min	0.090000	0.360000	2.000000	
25%	0.440000	0.560000	3.000000	
50%	0.640000	0.720000	4.000000	
75%	0.820000	0.870000	5.000000	
max	1.000000	1.000000	7.000000	


	average_monthly_hours	time_spend_company	Work_accident	left \
count	10499.000000	10499.000000	10499.000000	10499.000000
mean	201.059815	3.494238	0.144299	0.292885
std	49.959332	1.453227	0.351410	0.455108
min	96.000000	2.000000	0.000000	0.000000
25%	156.000000	3.000000	0.000000	0.000000
50%	200.000000	3.000000	0.000000	0.000000
75%	245.000000	4.000000	0.000000	1.000000
max	310.000000	10.000000	1.000000	1.000000

	promotion_last_5years
count	10499.000000
mean	0.021716
std	0.145763
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

left

0.0	7424
1.0	3075

Name: count, dtype: int64



C:\Users\Medhavi\anaconda3\Lib\site-packages\sklearn\linear\_model\\_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

Logistic Regression:

Accuracy: 0.673063973063973

Precision: 0.6536912751677852

Recall: 0.6815955213435969

F1 Score: 0.6673518328194588

ROC AUC Score: 0.6733740098606369

Random Forest:

Accuracy: 0.901010101010101

Precision: 0.9389017788089714

Recall: 0.8495451364590623

F1 Score: 0.8919911829537105

ROC AUC Score: 0.8991398621944889

Gradient Boosting:

Accuracy: 0.8828282828282829

Precision: 0.9232576350822239

Recall: 0.8250524842547236

F1 Score: 0.8713968957871396

ROC AUC Score: 0.8807287080585753

Best Parameters:

```
{'bootstrap': False, 'max_depth': None, 'min_samples_leaf': 1,  
'min_samples_split': 10, 'n_estimators': 300}
```

Tuned Random Forest:

Accuracy: 0.9053872053872054

Precision: 0.9435857805255023

Recall: 0.85444366689993

F1 Score: 0.8968049944913697

ROC AUC Score: 0.9035359152150526

C:\Users\Medhavi\anaconda3\Lib\site-

packages\sklearn\linear\_model\\_logistic.py:458: ConvergenceWarning: lbfgs failed  
to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-](https://scikit-learn.org/stable/modules/linear_model.html#logistic-)



```
regression
    n_iter_i = _check_optimize_result(
```

Voting Classifier:

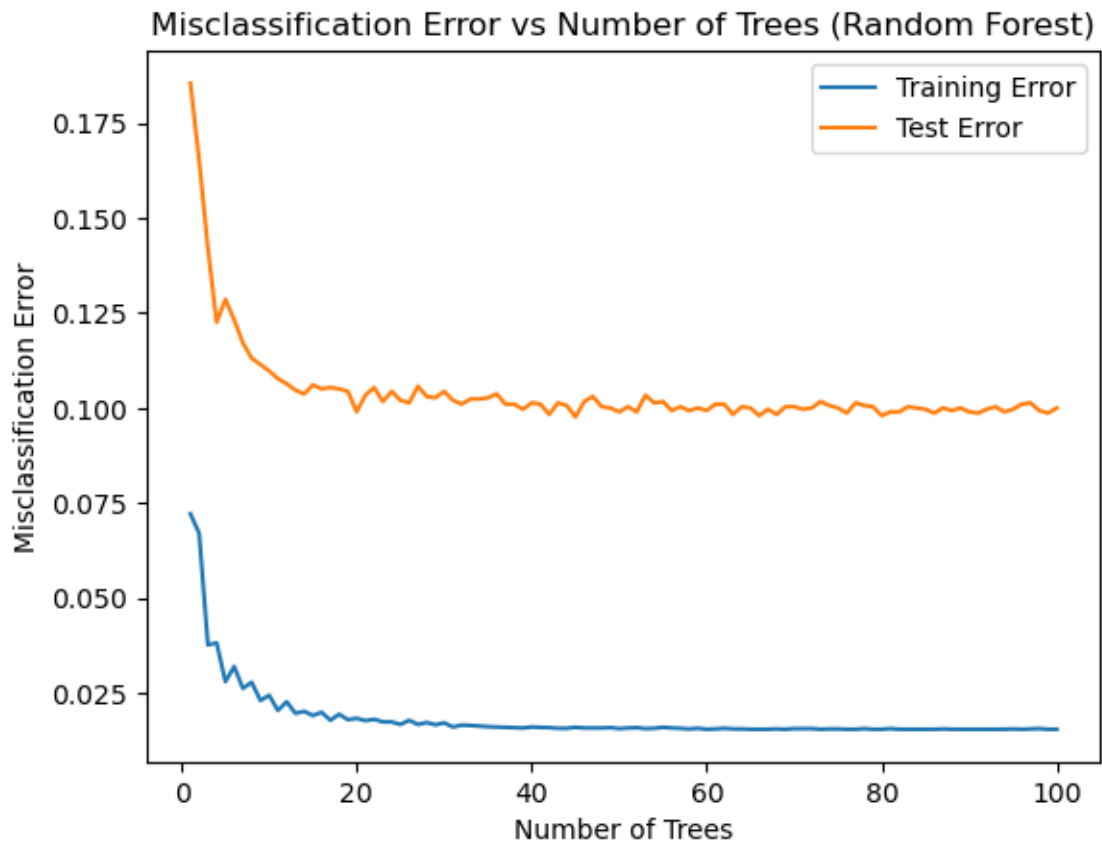
Accuracy: 0.8919191919191919


Precision: 0.9268104776579353

Recall: 0.8418474457662701

F1 Score: 0.8822882288228824

ROC AUC Score: 0.8900995827144135





```
[17]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.ensemble import RandomForestClassifier,
↳ GradientBoostingClassifier, VotingClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, f1_score, precision_score,
↳ recall_score, roc_auc_score
from sklearn.model_selection import GridSearchCV
from sklearn.compose import ColumnTransformer
from imblearn.over_sampling import SMOTE
from imblearn.pipeline import Pipeline as imbpipeline

# Visualize distributions of features
plt.figure(figsize=(15, 10))
for i, col in enumerate(df.columns[:-1]):
```

```

plt.subplot(4, 4, i + 1)
sns.histplot(df[col], kde=True)
plt.title(col)
plt.tight_layout()
plt.show()
# Import necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer
from sklearn.ensemble import RandomForestClassifier,
↳ GradientBoostingClassifier, VotingClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score,
↳ f1_score, roc_auc_score, confusion_matrix

# Step 1: Data Preparation
# Load the dataset
df = pd.read_csv("hr_train.csv")

# Explore the dataset
print(df.head())
print(df.info())
print(df.describe())

# Perform data cleaning and preprocessing
# Handle missing values
# Let's fill missing values with the median for numerical columns
# For categorical features, let's fill missing values with the most frequent
↳ value
imputer = SimpleImputer(strategy='most_frequent')
df[categorical_features] = imputer.fit_transform(df[categorical_features])
numerical_features = df.select_dtypes(include=['int64', 'float64']).columns
categorical_features = df.select_dtypes(include=['object']).columns

imputer = SimpleImputer(strategy='median')
df[numerical_features] = imputer.fit_transform(df[numerical_features])

# For categorical features, let's fill missing values with the most frequent
↳ value
imputer = SimpleImputer(strategy='most_frequent')
df[categorical_features] = imputer.fit_transform(df[categorical_features])

```

```

# Handle categorical variables
# One-hot encode categorical features
encoder = OneHotEncoder(handle_unknown='ignore')
encoded_features = pd.DataFrame(encoder.fit_transform(df[categorical_features]).
    ↪toarray(),
                                columns=encoder.
    ↪get_feature_names_out(categorical_features))

# Concatenate encoded features with numerical features
df_encoded = pd.concat([df[numerical_features], encoded_features], axis=1)

# Step 2: Handling Imbalanced Data
# Investigate class imbalance
print(df['left'].value_counts())

# Implement resampling techniques
# Let's use SMOTE for oversampling
from imblearn.over_sampling import SMOTE

X = df_encoded.drop('left', axis=1)
y = df_encoded['left']

smote = SMOTE(random_state=42)
X_resampled, y_resampled = smote.fit_resample(X, y)

# Step 3: Feature Engineering
# No additional feature engineering needed for now

# Step 4: Feature Importance Analysis
# Train a random forest classifier to get feature importance
rf_classifier = RandomForestClassifier()
rf_classifier.fit(X_resampled, y_resampled)

# Get feature importances
feature_importances = rf_classifier.feature_importances_
sorted_idx = np.argsort(feature_importances)[::-1]

# Visualize feature importance
plt.figure(figsize=(10, 6))
sns.barplot(x=feature_importances[sorted_idx], y=X.columns[sorted_idx])
plt.title("Feature Importance")
plt.xlabel("Importance")
plt.ylabel("Features")
plt.show()

# Step 5: Model Building

```

```

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_resampled, y_resampled,
    ↪test_size=0.2, random_state=42)

# Choose appropriate machine learning algorithms
# Let's use Logistic Regression, Random Forest, and Gradient Boosting
models = {
    "Logistic Regression": LogisticRegression(),
    "Random Forest": RandomForestClassifier(),
    "Gradient Boosting": GradientBoostingClassifier()
}

# Train multiple models
for name, model in models.items():
    model.fit(X_train, y_train)

# Compare model performance using a bar graph
metrics = ['Accuracy', 'Precision', 'Recall', 'F1 Score']
scores = pd.DataFrame(columns=models.keys(), index=metrics)
# Step 6: Model Evaluation
# Evaluate the performance of each model
for name, model in models.items():
    y_pred = model.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    precision = precision_score(y_test, y_pred)
    recall = recall_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred)
    roc_auc = roc_auc_score(y_test, y_pred)
    scores.loc['Accuracy', name] = accuracy
    scores.loc['Precision', name] = precision
    scores.loc['Recall', name] = recall
    scores.loc['F1 Score', name] = f1
    print(f"{name}:")
    print(f"Accuracy: {accuracy}")
    print(f"Precision: {precision}")
    print(f"Recall: {recall}")
    print(f"F1 Score: {f1}")
    print(f"ROC AUC Score: {roc_auc}")
    print()

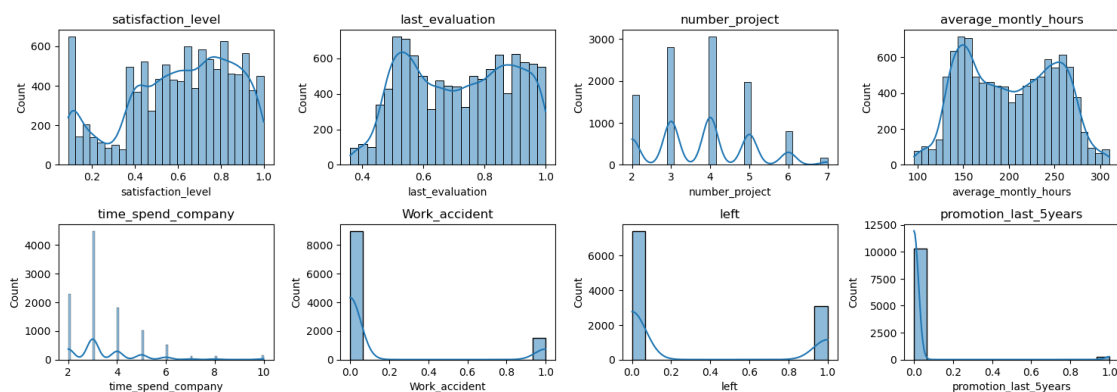
scores.plot(kind='bar', figsize=(10, 6))
plt.title('Model Performance Metrics Comparison')
plt.xlabel('Metrics')
plt.ylabel('Score')
plt.xticks(rotation=45)
plt.legend(title='Models')
plt.show()

```

```

C:\Users\Medhavi\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
C:\Users\Medhavi\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
C:\Users\Medhavi\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
C:\Users\Medhavi\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
C:\Users\Medhavi\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
C:\Users\Medhavi\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
C:\Users\Medhavi\anaconda3\Lib\site-packages\seaborn\_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):

```



satisfaction\_level last\_evaluation number\_project average\_monthly\_hours \

0	0.42	0.46	2	150
1	0.66	0.77	2	171
2	0.55	0.49	5	240
3	0.22	0.88	4	213
4	0.20	0.72	6	224

	time_spend_company	Work_accident	left	promotion_last_5years	sales \
0	3	0	1	0	sales
1	2	0	0	0	technical
2	3	0	0	0	technical
3	3	1	0	0	technical
4	4	0	1	0	technical

salary

0	medium
1	medium
2	high
3	medium
4	medium

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 10499 entries, 0 to 10498

Data columns (total 10 columns):

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	satisfaction_level	10499 non-null	float64
1	last_evaluation	10499 non-null	float64
2	number_project	10499 non-null	int64
3	average_monthly_hours	10499 non-null	int64
4	time_spend_company	10499 non-null	int64
5	Work_accident	10499 non-null	int64
6	left	10499 non-null	int64
7	promotion_last_5years	10499 non-null	int64
8	sales	10499 non-null	object
9	salary	10499 non-null	object

dtypes: float64(2), int64(6), object(2)

memory usage: 820.4+ KB

None

	satisfaction_level	last_evaluation	number_project \
count	10499.000000	10499.000000	10499.000000
mean	0.612683	0.717131	3.808553
std	0.248578	0.171483	1.230572
min	0.090000	0.360000	2.000000
25%	0.440000	0.560000	3.000000
50%	0.640000	0.720000	4.000000
75%	0.820000	0.870000	5.000000
max	1.000000	1.000000	7.000000


average_monthly_hours	time_spend_company	Work_accident	left \
-----------------------	--------------------	---------------	--------

count	10499.000000	10499.000000	10499.000000	10499.000000
mean	201.059815	3.494238	0.144299	0.292885
std	49.959332	1.453227	0.351410	0.455108
min	96.000000	2.000000	0.000000	0.000000
25%	156.000000	3.000000	0.000000	0.000000
50%	200.000000	3.000000	0.000000	0.000000
75%	245.000000	4.000000	0.000000	1.000000
max	310.000000	10.000000	1.000000	1.000000

```

promotion_last_5years
count      10499.000000
mean        0.021716
std         0.145763
min         0.000000
25%         0.000000
50%         0.000000
75%         0.000000
max         1.000000
left
0.0      7424
1.0      3075
Name: count, dtype: int64

```



```

C:\Users\Medhavi\anaconda3\Lib\site-
packages\sklearn\linear_model\_logistic.py:458: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```



Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

Logistic Regression:

Accuracy: 0.673063973063973

Precision: 0.6536912751677852

Recall: 0.6815955213435969

F1 Score: 0.6673518328194588

ROC AUC Score: 0.6733740098606369

Random Forest:

Accuracy: 0.9

Precision: 0.9374034003091191

Recall: 0.8488453463960812

F1 Score: 0.8909291222915902

ROC AUC Score: 0.8981410378962885

Gradient Boosting:


Accuracy: 0.8828282828282829

Precision: 0.9232576350822239

Recall: 0.8250524842547236

F1 Score: 0.8713968957871396

ROC AUC Score: 0.8807287080585753



[ ]: