

```
# `pip3 install assemblyai` (macOS)
# `pip install assemblyai` (Windows)

import assemblyai as aai

aai.settings.api_key = "dfc5fb97b35e4fde9f97900564047539"
transcriber = aai.Transcriber()

transcript =
transcriber.transcribe("https://storage.googleapis.com/aai-web-
samples/news.mp4")
# transcript = transcriber.transcribe("./my-local-audio-file.wav")

print(transcript.text)
```

I'm David Curley at the Smithsonian Air and Space Museum, where we are marking 50 years since man landed and walked on the moon in a lander just like this one. We are going to show you some of the actual ABC News coverage from 50 years ago. During that eight day mission of this remarkable achievement, Apollo Eleven's lander, the Eagle, would be the first man craft to land on the moon. For training, NASA came up with an unusual contraption. Neil Armstrong actually had to eject from it once, and then he had a couple of successful flights. ABC News anchor at the time, 50 years ago, Frank Reynolds with a look at that unusual trainer. Apollo Eleven Commander Neil Armstrong is at the controls of a lunar landing training vehicle, testing the reaction control jets. These thrusters stabilize the LEM during landing and takeoff. The LLTV is designed to simulate the behavior of the LEM as it lands in the moon's gravity. Lunar gravity is one 6th that of the Earth's. Neil Armstrong flew one of these vehicles on May 6, 1968, and that flight was nearly his last. Later reports by a NASA investigating team said the crash, which we'll see soon, was caused by a loss of fuel pressure compounded by a warning light that failed to work. Armstrong's coolness under pressure saved himself and possibly months of delay for the Apollo program. Later that same year, 1968, another test pilot, Joseph Allegranti, also escaped from an LLTV just before it crashed. The training vehicle then underwent several design modifications and improvements. Engineers had to increase the vehicle's rocket power to help stabilize the craft. That made the LLTV less of a moon gravity simulator, but it improved pilot safety. On June 16, 1969, Neil Armstrong flew the vehicle, sometimes called the flying bedstead, to several perfect landings. The question was, how does the machine fly? And the answer is that we're very pleased with the way it flies. It's a significant improvement over the LLRV, which we were flying here a year ago, and I think it does an excellent job of actually capturing the handling characteristics of the lunar module in landing maneuver. It's really a great deal different than any other kind of aircraft that I've ever flown. The simulation of lunar gravity has some aspects that make this type of flight sufficiently different from anything else we've ever done to make this vehicle very

worthwhile. And I'm very pleased that I had the opportunity to get some flights in it here just before the Apollo eleven flight.

```
import tensorflow as tf
import os
from PIL import Image
import numpy as np

def preprocess_image(image):
    image = tf.image.resize(image, (64, 64))
    image = (image - 127.5) / 127.5 # Normalize to [-1, 1]
    return image

def load_data(data_dir):
    images = []
    for img_file in os.listdir(data_dir):
        img_path = os.path.join(data_dir, img_file)
        img = Image.open(img_path).convert('RGB')
        img = np.array(img)
        img = preprocess_image(img)
        images.append(img)
    dataset = tf.data.Dataset.from_tensor_slices(images).batch(128)
    return dataset

data_dir = 'CelebA-HQ-img'
dataset = load_data(data_dir)

def make_generator_model():
    model = tf.keras.Sequential()
    model.add(tf.keras.layers.Dense(8*8*256, use_bias=False,
input_shape=(100,)))
    model.add(tf.keras.layers.BatchNormalization())
    model.add(tf.keras.layers.LeakyReLU())
    model.add(tf.keras.layers.Reshape((8, 8, 256)))
    model.add(tf.keras.layers.Conv2DTranspose(128, (5, 5), strides=(2,
2), padding='same', use_bias=False))
    model.add(tf.keras.layers.BatchNormalization())
    model.add(tf.keras.layers.LeakyReLU())
    model.add(tf.keras.layers.Conv2DTranspose(64, (5, 5), strides=(2,
2), padding='same', use_bias=False))
    model.add(tf.keras.layers.BatchNormalization())
    model.add(tf.keras.layers.LeakyReLU())
    model.add(tf.keras.layers.Conv2DTranspose(3, (5, 5), strides=(2,
2), padding='same', use_bias=False, activation='tanh'))
    return model

def make_discriminator_model():
    model = tf.keras.Sequential()
    model.add(tf.keras.layers.Conv2D(64, (5, 5), strides=(2, 2),
padding='same', input_shape=[64, 64, 3]))
```

```

        model.add(tf.keras.layers.LeakyReLU())
        model.add(tf.keras.layers.Dropout(0.3))
        model.add(tf.keras.layers.Conv2D(128, (5, 5), strides=(2, 2),
padding='same'))
        model.add(tf.keras.layers.LeakyReLU())
        model.add(tf.keras.layers.Dropout(0.3))
        model.add(tf.keras.layers.Flatten())
        model.add(tf.keras.layers.Dense(1))
        return model

generator = make_generator_model()
discriminator = make_discriminator_model()

cross_entropy = tf.keras.losses.BinaryCrossentropy(from_logits=True)

def generator_loss(fake_output):
    return cross_entropy(tf.ones_like(fake_output), fake_output)

def discriminator_loss(real_output, fake_output):
    real_loss = cross_entropy(tf.ones_like(real_output), real_output)
    fake_loss = cross_entropy(tf.zeros_like(fake_output), fake_output)
    return real_loss + fake_loss

generator_optimizer = tf.keras.optimizers.Adam(1e-4)
discriminator_optimizer = tf.keras.optimizers.Adam(1e-4)

import os
import matplotlib.pyplot as plt

EPOCHS = 50
noise_dim = 100
num_examples_to_generate = 16

seed = tf.random.normal([num_examples_to_generate, noise_dim])

@tf.function
def train_step(images):
    noise = tf.random.normal([128, noise_dim])

    with tf.GradientTape() as gen_tape, tf.GradientTape() as
disc_tape:
        generated_images = generator(noise, training=True)

        real_output = discriminator(images, training=True)
        fake_output = discriminator(generated_images, training=True)

        gen_loss = generator_loss(fake_output)
        disc_loss = discriminator_loss(real_output, fake_output)

    gradients_of_generator = gen_tape.gradient(gen_loss,

```

```

generator.trainable_variables)
    gradients_of_discriminator = disc_tape.gradient(disc_loss,
discriminator.trainable_variables)

    generator_optimizer.apply_gradients(zip(gradients_of_generator,
generator.trainable_variables))

discriminator_optimizer.apply_gradients(zip(gradients_of_discriminator
, discriminator.trainable_variables))

def train(dataset, epochs):
    for epoch in range(epochs):
        for image_batch in dataset:
            train_step(image_batch)
        print(f'Epoch {epoch+1} completed')
        generate_and_save_images(generator, epoch + 1, seed)

def generate_and_save_images(model, epoch, test_input):
    predictions = model(test_input, training=False)
    fig = plt.figure(figsize=(4, 4))
    for i in range(predictions.shape[0]):
        plt.subplot(4, 4, i + 1)
        plt.imshow((predictions[i] * 127.5 +
127.5).numpy().astype("uint8"))
        plt.axis('off')
    plt.savefig(f'image_at_epoch_{epoch:04d}.png')
    plt.show()

```

```
train(dataset, EPOCHS)
```


```
C:\Users\Medhavi\anaconda3\Lib\site-packages\keras\src\layers\core\
dense.py:87: UserWarning: Do not pass an `input_shape`/`input_dim`
argument to a layer. When using Sequential models, prefer using an
`Input(shape)` object as the first layer in the model instead.
```

```
    super().__init__(activity_regularizer=activity_regularizer,
**kwargs)
```


```
C:\Users\Medhavi\anaconda3\Lib\site-packages\keras\src\layers\
convolutional\base_conv.py:107: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential
models, prefer using an `Input(shape)` object as the first layer in
the model instead.
```

```
    super().__init__(activity_regularizer=activity_regularizer,
**kwargs)
```

```
Epoch 1 completed
```




Epoch 2 completed



Epoch 3 completed



Epoch 4 completed




Epoch 5 completed




Epoch 6 completed



Epoch 7 completed




Epoch 8 completed



Epoch 9 completed



Epoch 10 completed




Epoch 11 completed



Epoch 12 completed



Epoch 13 completed



Epoch 14 completed



Epoch 15 completed



Epoch 16 completed



Epoch 17 completed




Epoch 18 completed



Epoch 19 completed




Epoch 20 completed




Epoch 21 completed



Epoch 22 completed



Epoch 23 completed




Epoch 24 completed



Epoch 25 completed



Epoch 26 completed



Epoch 27 completed



Epoch 28 completed



Epoch 29 completed



Epoch 30 completed



Epoch 31 completed



Epoch 32 completed



Epoch 33 completed



Epoch 34 completed



Epoch 35 completed



Epoch 36 completed



Epoch 37 completed



Epoch 38 completed



Epoch 39 completed




Epoch 40 completed



Epoch 41 completed




Epoch 42 completed



Epoch 43 completed




Epoch 44 completed



Epoch 45 completed



Epoch 46 completed



Epoch 47 completed



Epoch 48 completed



Epoch 49 completed



Epoch 50 completed



```
import tensorflow as tf
from tensorflow.keras import layers
import tensorflow_datasets as tfds

# Load and preprocess the CelebA dataset
def preprocess_image(image, label):
    image = tf.image.resize(image, [64, 64])
    image = (image - 127.5) / 127.5 # Normalize to [-1, 1]
```

```

    return image
def load_data(data_dir):
    images = []
    for img_file in os.listdir(data_dir):
        img_path = os.path.join(data_dir, img_file)
        img = Image.open(img_path).convert('RGB')
        img = np.array(img)
        img = preprocess_image(img)
        images.append(img)
    dataset = tf.data.Dataset.from_tensor_slices(images).batch(128)
    return dataset

```

```

data_dir = 'CelebA-HQ-img'
dataset = load_data(data_dir)

```

Define the generator model

```

def make_generator_model():
    model = tf.keras.Sequential()
    model.add(layers.Dense(8*8*256, use_bias=False,
input_shape=(100,)))
    model.add(layers.BatchNormalization())
    model.add(layers.LeakyReLU())

    model.add(layers.Reshape((8, 8, 256)))
    model.add(layers.Conv2DTranspose(128, (5, 5), strides=(2, 2),
padding='same', use_bias=False))
    model.add(layers.BatchNormalization())
    model.add(layers.LeakyReLU())

    model.add(layers.Conv2DTranspose(64, (5, 5), strides=(2, 2),
padding='same', use_bias=False))
    model.add(layers.BatchNormalization())
    model.add(layers.LeakyReLU())

    model.add(layers.Conv2DTranspose(3, (5, 5), strides=(2, 2),
padding='same', use_bias=False, activation='tanh'))

    return model

```

Define the discriminator model

```

def make_discriminator_model():
    model = tf.keras.Sequential()
    model.add(layers.Conv2D(64, (5, 5), strides=(2, 2),
padding='same', input_shape=[64, 64, 3]))
    model.add(layers.LeakyReLU())
    model.add(layers.Dropout(0.3))

    model.add(layers.Conv2D(128, (5, 5), strides=(2, 2),
padding='same'))

```

```

model.add(layers.LeakyReLU())
model.add(layers.Dropout(0.3))

model.add(layers.Flatten())
model.add(layers.Dense(1))

return model

# Define loss and optimizers
cross_entropy = tf.keras.losses.BinaryCrossentropy(from_logits=True)

def discriminator_loss(real_output, fake_output):
    real_loss = cross_entropy(tf.ones_like(real_output), real_output)
    fake_loss = cross_entropy(tf.zeros_like(fake_output), fake_output)
    total_loss = real_loss + fake_loss
    return total_loss

def generator_loss(fake_output):
    return cross_entropy(tf.ones_like(fake_output), fake_output)

generator = make_generator_model()
discriminator = make_discriminator_model()

generator_optimizer = tf.keras.optimizers.Adam(1e-4)
discriminator_optimizer = tf.keras.optimizers.Adam(1e-4)

# Training loop
@tf.function
def train_step(images):
    noise = tf.random.normal([256, 100])

    with tf.GradientTape() as gen_tape, tf.GradientTape() as disc_tape:
        generated_images = generator(noise, training=True)

        real_output = discriminator(images, training=True)
        fake_output = discriminator(generated_images, training=True)

        gen_loss = generator_loss(fake_output)
        disc_loss = discriminator_loss(real_output, fake_output)

        gradients_of_generator = gen_tape.gradient(gen_loss,
generator.trainable_variables)
        gradients_of_discriminator = disc_tape.gradient(disc_loss,
discriminator.trainable_variables)

        generator_optimizer.apply_gradients(zip(gradients_of_generator,
generator.trainable_variables))

    discriminator_optimizer.apply_gradients(zip(gradients_of_discriminator
, discriminator.trainable_variables))

```

```

def train(dataset, epochs):
    for epoch in range(epochs):
        for image_batch in dataset:
            train_step(image_batch)

        # Produce images for the GIF as we go
        display.clear_output(wait=True)
        generate_and_save_images(generator, epoch + 1, seed)

    # Generate after the final epoch
    display.clear_output(wait=True)
    generate_and_save_images(generator, epochs, seed)

def generate_and_save_images(model, epoch, test_input):
    predictions = model(test_input, training=False)
    fig = plt.figure(figsize=(4, 4))

    for i in range(predictions.shape[0]):
        plt.subplot(4, 4, i + 1)
        plt.imshow((predictions[i] * 127.5 +
127.5).numpy().astype("uint8"))
        plt.axis('off')

    plt.savefig('image_at_epoch_{:04d}.png'.format(epoch))
    plt.show()

# Load the dataset
dataset = load_data()

# Train the model
train(dataset, epochs=50)

```

```

-----
-----
TypeError                                Traceback (most recent call
last)

```

```

Cell In[6], line 22
    19     return dataset
    21 data_dir = 'CelebA-HQ-img'
--> 22 dataset = load_data(data_dir)
    25 # Define the generator model
    26 def make_generator_model():

```

```

Cell In[6], line 16, in load_data(data_dir)
    14     img = Image.open(img_path).convert('RGB')
    15     img = np.array(img)
--> 16     img = preprocess_image(img)
    17     images.append(img)
    18 dataset =

```



```
tf.data.Dataset.from_tensor_slices(images).batch(128)
```

```
TypeError: preprocess_image() missing 1 required positional argument: 'label'
```

```
import torch
import torchvision.transforms as transforms
from torchvision.models import inception_v3
import numpy as np
from scipy.linalg import sqrtm
from torch.utils.data import DataLoader, Dataset
from PIL import Image
import os
from tqdm import tqdm

class ImageDataset(Dataset):
    def __init__(self, image_paths, transform=None):
        self.image_paths = image_paths
        self.transform = transform

    def __len__(self):
        return len(self.image_paths)

    def __getitem__(self, idx):
        image = Image.open(self.image_paths[idx]).convert("RGB")
        if self.transform:
            image = self.transform(image)
        return image

def load_images(image_folder):
    image_paths = [os.path.join(image_folder, img) for img in
os.listdir(image_folder)]
    transform = transforms.Compose([
        transforms.Resize((299, 299)),
        transforms.ToTensor(),
        transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229,
0.224, 0.225]),
    ])
    dataset = ImageDataset(image_paths, transform=transform)
    return DataLoader(dataset, batch_size=32, shuffle=False)

def get_inception_activations(dataloader, model, device):
    model.eval()
    activations = []
    with torch.no_grad():
        for batch in tqdm(dataloader):
            batch = batch.to(device)
            pred = model(batch)[0]
            activations.append(pred.cpu().numpy())
    activations = np.concatenate(activations, axis=0)
```

```

    return activations

def calculate_inception_score(activations, splits=10):
    scores = []
    for i in range(splits):
        part = activations[i * (activations.shape[0] // splits): (i +
1) * (activations.shape[0] // splits), :]
        py = np.mean(part, axis=0)
        scores.append(np.exp(np.mean(np.sum(part * (np.log(part + 1e-
10) - np.log(py + 1e-10)), axis=1))))
    return np.mean(scores), np.std(scores)

def calculate_fid(real_activations, fake_activations):
    mu1, sigma1 = np.mean(real_activations, axis=0),
np.cov(real_activations, rowvar=False)
    mu2, sigma2 = np.mean(fake_activations, axis=0),
np.cov(fake_activations, rowvar=False)
    ssdiff = np.sum((mu1 - mu2)**2.0)
    covmean = sqrtm(sigma1.dot(sigma2))
    if np.iscomplexobj(covmean):
        covmean = covmean.real
    return ssdiff + np.trace(sigma1 + sigma2 - 2.0 * covmean)

def main(real_images_folder, generated_images_folder):
    device = torch.device("cuda" if torch.cuda.is_available() else
"cpu")
    model = inception_v3(pretrained=True,
transform_input=False).to(device)

    real_dataloader = load_images(real_images_folder)
    generated_dataloader = load_images(generated_images_folder)

    real_activations = get_inception_activations(real_dataloader,
model, device)
    fake_activations = get_inception_activations(generated_dataloader,
model, device)

    inception_score_mean, inception_score_std =
calculate_inception_score(fake_activations)
    fid_score = calculate_fid(real_activations, fake_activations)

    print(f"Inception Score: {inception_score_mean} ±
{inception_score_std}")
    print(f"FID Score: {fid_score}")

if __name__ == "__main__":
    real_images_folder = "CelebA-HQ-img"
    generated_images_folder = "generated_images"
    main(real_images_folder, generated_images_folder)

```

```
C:\Users\Medhavi\anaconda3\Lib\site-packages\torchvision\models\_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the future, please use 'weights' instead.
```

```
warnings.warn(
```

```
C:\Users\Medhavi\anaconda3\Lib\site-packages\torchvision\models\
_utils.py:223: UserWarning: Arguments other than a weight enum or
`None` for 'weights' are deprecated since 0.13 and may be removed in
the future. The current behavior is equivalent to passing
`weights=Inception_V3_Weights.IMAGENET1K_V1`. You can also use
`weights=Inception_V3_Weights.DEFAULT` to get the most up-to-date
weights.
```

```
warnings.warn(msg)
```

100% |

```
| 938/938 [52:22<00:00, 3.35s/it]
```

100% |

```
| 2/2 [00:03<00:00, 1.60s/it]
```

[illegible]

Cell In[1], line 82

```
80 real_images_folder = "CelebA-HQ-img"
```

```
81 generated_images_folder = "generated_images"
```

```
---> 82 main(real_images_folder, generated_images_folder)
```

```
Cell In[1], line 73, in main(real_images_folder,
generated_images_folder)
```

```
70 real_activations = get_inception_activations(real_dataloader,  
model, device)
```

```
71 fake_activations =
```

```
get_inception_activations(generated_dataloader, model, device)
```

```
---> 73 inception score mean, inception score std =
```

```
calculate_inception_score(fake_activations)
```

```
74 fid_score = calculate_fid(real_activations, fake_activations)
```

```
76 print(f"Inception Score: {inception_score_mean} ±
```

```
{inception score std}")
```

Cell In[1], line 49, in calculate_inception_score(activations, splits)

```
47 scores = []
```

```
48 for i in range(splits):
```

```

--> 49     part = activations[i * (activations.shape[0] // splits):
(i + 1) * (activations.shape[0] // splits), :]

```

```
50     py = np.mean(part, axis=0)
```

```
51     scores.append(np.exp(np.mean(np.sum(part * (np.log(part +
1e-10) - np.log(py + 1e-10)), axis=1))))
```

IndexError: too many indices for array: array is 1-dimensional, but 2 were indexed

```
import os
import cv2
import numpy as np
from skimage.metrics import structural_similarity as ssim

def calculate_psnr(img1, img2):
    mse = np.mean((img1 - img2) ** 2)
    if mse == 0:
        return 100
    PIXEL_MAX = 255.0
    return 20 * np.log10(PIXEL_MAX / np.sqrt(mse))

def calculate_ssim(img1, img2):
    return ssim(img1, img2, multichannel=True)

def load_images_from_folder(folder):
    images = []
    for filename in os.listdir(folder):
        img_path = os.path.join(folder, filename)
        img = cv2.imread(img_path)
        if img is not None:
            images.append((filename, img))
        else:
            print(f"Warning: {filename} could not be loaded.")
    return images

def main():
    folder = 'generated_images'
    reference_image_path = 'generated_images/50.png'

    reference_image = cv2.imread(reference_image_path)
    if reference_image is None:
        print(f"Error: Reference image at path {reference_image_path}
could not be loaded.")
        return

    images = load_images_from_folder(folder)

    results = []

    for (filename, image) in images:
        if image.shape != reference_image.shape:
            print(f"Warning: {filename} has a different shape than the
reference image. Skipping.")
            continue
        psnr_value = calculate_psnr(reference_image, image)
        ssim_value = calculate_ssim(reference_image, image)
```



```
        results.append((filename, psnr_value, ssim_value))

    for result in results:
        print(f"Filename: {result[0]}, PSNR: {result[1]:.2f}, SSIM:
{result[2]:.4f}")

if __name__ == "__main__":
    main()
```

Error: Reference image at path generated_images/50.png could not be loaded.